W207 Final Project

AMES Housing

Mrinal Chawla, Thomas Gao, Fengyao Luo

The Inference Problem

GIVEN
● A vector of features about the house ○ Neighborhood ○ Quality ■ Overall, Pool, etc. ○ Sale Condition ■ Normal, Abnormal ○ Amenities ■ Alley, Garage, Basement, etc. ○ Size ■ Sqft, # of rooms, # of bathrooms etc

PREDICT
● Sale Price of the House

WHY
● Helps housing market for both sellers and buyers ● Help plan renovations ● Help plan infrastructure improvements ● Etc.

Exploratory Data Analysis

- SalePrice slightly right-skewed
  - Log transformation to fix
- Sizing variables highly correlated
  - Bsmt sqft vs 1st floor sqft
  - GarageArea vs GarageCars

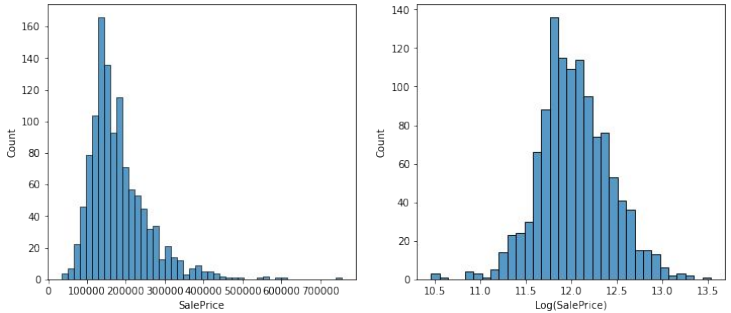

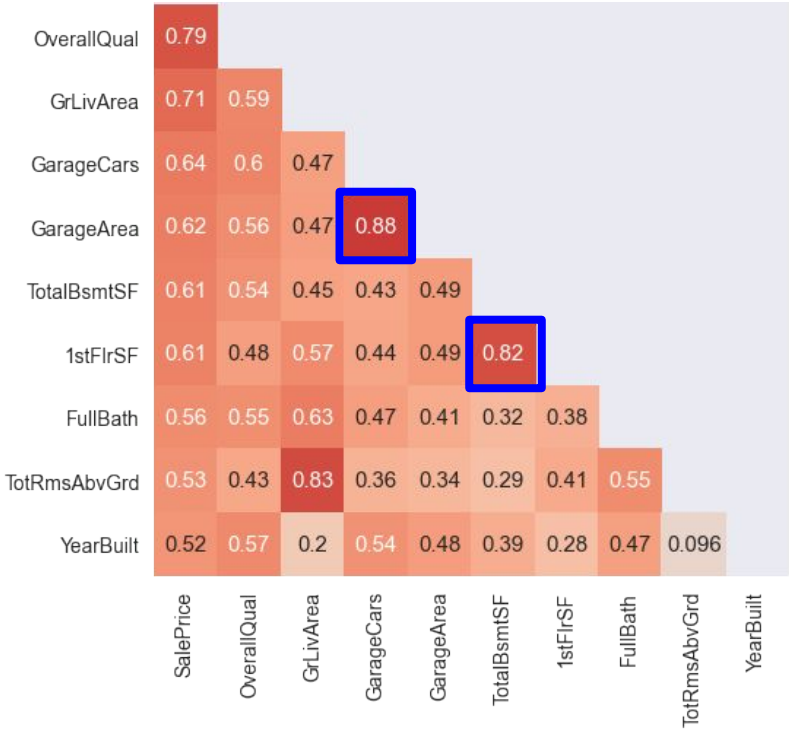

Exploratory Data Analysis

- Generally linear relationships
- Few outliers for expensive homes
- Top indicators
  - Overall Quality
  - Living Area
  - Neighborhood
- Scaling and encoding

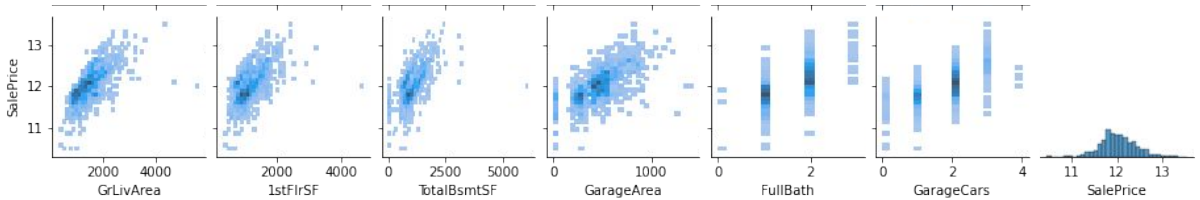

Baseline Models

PREDICT MEAN

Predict average sale price for every house

RMSE: 0.419

LINEAR REGRESSION

Use top 2 features and neighborhood

RMSE: 0.169

LASSO

Linear Regression with L1 Regularization

RMSE: 0.137

Error Analysis

- Complexity vs Performance
  - Number of Features
  - RMSE
- Outliers
  - Living Area
  - Quality
- Non-Linear Relationship
  - YrSold
  - SalePrice

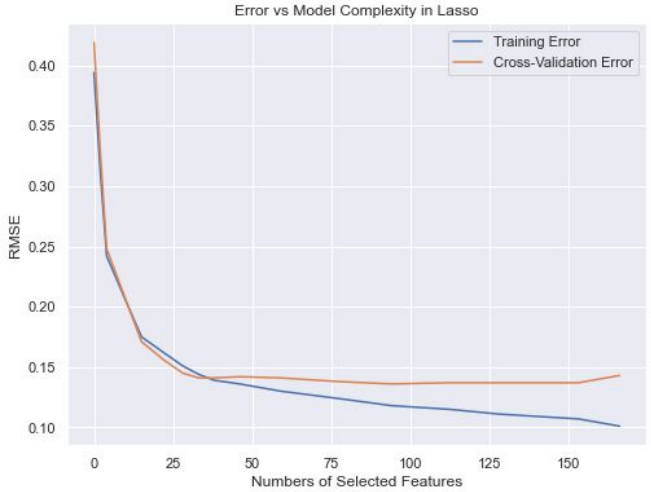

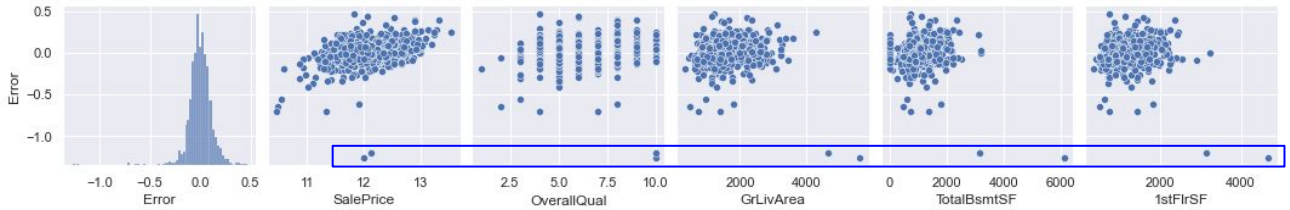

Feature Engineering

- Aggregate size features
  - Total Sqft
  - Average Room Sqft
  - Total bathrooms
  - Total porch sqft
- Presence of amenities (binary)
  - Alley
  - Garage
  - Basement
  - Pool
- Years Since Remodelled (Year Sold - Year Remodelled)
- Seasonality (Month Sold → Season)
- Skewness (np.log)
- Neighborhood bins

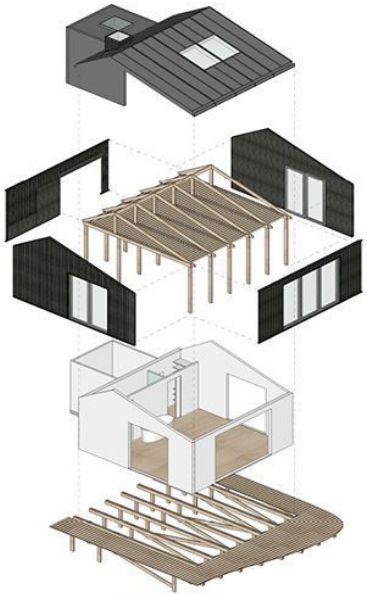

Final Models

LASSO	SPLIT LASSO
Linear Regression with L1 Regularization RMSE: 0.121	Linear Regression with L1 Regularization + Separate Model for each Neighborhood Bin RMSE: 0.126
RANDOM FOREST	ENSEMBLE
Nonlinear Model RMSE: 0.135	0.75 Lasso + 0.25 Random Forest RMSE: 0.119

Final Performance

- Model: Ensemble
- RMSE: 0.1217
- Kaggle: 10th percentile